

Poll: Why is multiplying numbers by 100 easier than multiplying by 128?

EECS 370 - Lecture 2

Binary and
Instruction Set Architecture (ISA)



Live Poll + Q&A: [slido.com #eecs370](https://slido.com/#eecs370) Poll and Q&A Link

Announcements


- P1a
 - Posted today
- Labs
 - No lab Monday (holiday)
 - Lab 1 due next Wednesday
 - Attendance starts next Friday
- OH
 - Started

My Office Hours

- 2 types:
- Group:
 - In-person
 - 30 minutes right after class (2901 BBB)
 - Prioritize group questions over individual debugging
 - Starting today
- Individual:
 - Some in-person, some virtual
 - See Google calendar for details
 - One-on-one: any questions welcome

Extra Resources

- Want more examples on binary? Two's complement?
 - See "resources tab" on website
 - Extra videos, review sheets



Instruction Set Architecture (ISA) Design Lectures


- **Lecture 2: ISA - storage types, binary and addressing modes**
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- Lecture 6 : Converting C to assembly – functions
- Lecture 7 : Translation software; libraries, memory layout

Agenda

- **Computer Model and Binary**
- ISAs
 - Registers
 - Control Flow
 - Representing Different Values

Basic Computer Model

- You know from 280 that computers have "memory"
 - Abstractly, a long array that holds values
- Every piece of data in a running program lives at a numerical **address** in memory
 - You can see the address in C by using the "&" operator



- Most programs work by loading values from memory to the processor, operating on those values, and writing values back into memory

Basic Memory Model

- 1st question in understanding how programs run on computers:
 - How are values actually represented in memory?
- Answer: binary

Aside: Decimal and Binary



- Humans represent numbers in base-10 (decimal) because we have 10 fingers (or "digits")
- The n^{th} digit corresponds to 10^n

$$1407 = 1 \cdot 10^3 + 4 \cdot 10^2 + 0 \cdot 10^1 + 7 \cdot 10^0 = 1000 + 400 + 00 + 7$$

Collection of 8 bits is called a byte

- Computers are made of wires with either high or low voltages
- Internally represents values in base-2 (binary) since it has "binary digits"
 - (or bits for short)
- In binary, the n^{th} bit corresponds to 2^n

$$1101 = 1 \cdot 2^3 + 1 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 = 8 + 4 + 0 + 1 = 13$$



Does Bart Simpson count in octal?



9

Aside: Hexadecimal

- A bunch of 0s and 1s is hard to read for humans
 - But translating to decimal and back is tricky
- Solution: Bases that are a power of 2 are easy to translate between, since a fixed group of bits corresponds to one digit
- In practice, base-16 or **hexadecimal** is used
 - Digits 0-9, plus letters A-F to represent 10-16



10

Aside: Hexadecimal

Represent binary using 0b. Hex using 0x. If not specified, it's decimal

- Every 4 bits corresponds to 1 hex digit (since $2^4=16$)

8 bits = 1 byte

(binary) 0b 0010 0101 1010 1011

(hexadecimal) 0x 2 5 A B

0x25AB 2 bytes



Live Poll + Q&A: slido.com #eecs370

11

Other Units in this Class

1 byte = 8 bits

Unit	Number of Bytes
word	4 (in this class)
Kilobyte (KB)	$2^{10} = 1,024$
Megabyte (MB)	$2^{20} = 1,048,576$
Gigabyte (GB)	$2^{30} = \text{About a billion}$



12

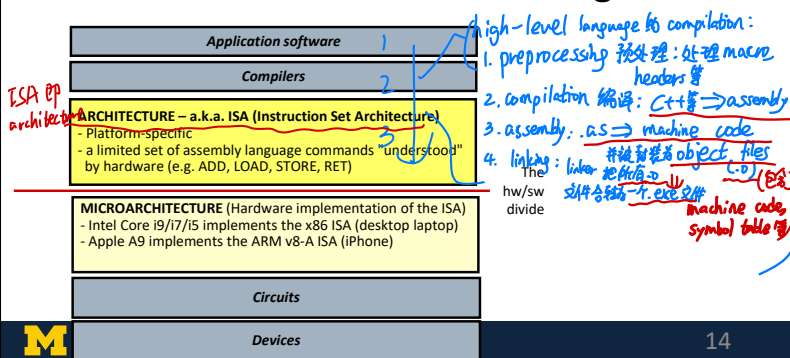
Agenda

- Computer Model and Binary
- ISAs
 - Registers
 - Control Flow
 - Representing Different Values



13

Where do ISAs come into the game?



14

Instruction Set Architecture (ISA): 指令集架构

- 包含:
- ① basic data types
 - ② instruction set
 - ③ register
 - ④ addressing mode
 - ⑤ ... 等等

ISA 是应用和硬件之间的桥梁

software app + compiler $\xrightarrow{\text{编译}}$ ISA + microarchitecture (implementation of circuits)

应用编译成使用 ISA 的 instructions, instructions 再 通过 microarchitecture 的 implementation 来控制硬件 (一种 microarch: 可以兼容不同的 ISA)

How is Assembly Different from C/C++?

- C/C++ instructions operate on **variables**
 - e.g. $x = i + j;$
 - Practically unlimited
- We might guess that assembly instructions act on addresses, e.g. $0x10000100 = 0x10000200 + 0x10000300$
- Problems:
 - This makes the instructions really long
 - As we'll see later in the course, memory is slow
 - We don't want to go multiple times for every instruction



15

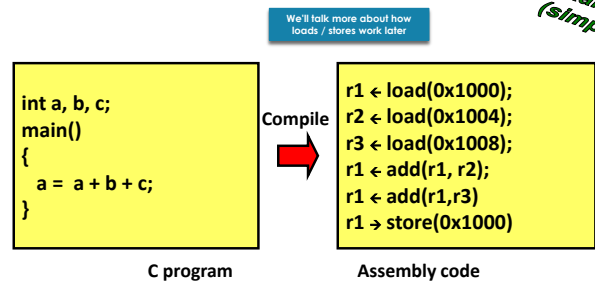
How is Assembly Different from C/C++?

- Modern ISAs define **registers**
 - Basically a small number (~8-32) of fixed-length, hardware variables that have simple names like "r5"
- In a **load-store architecture** (what we'll assume in this class):
 - load** instructions **bring** values from memory into a register
 - Other instructions specify register indices (compact and fast)
 - store** instructions send them back to memory



16

Example Assembly Code



17

Example Architectures

- ARMv8—LEGv8 subset from P+H text book
 - 32 registers (X0—X31) *(each 64 bits = 8 bytes)*
 - 64 bits in each register
 - Some have special uses e.g. X31 is always 0—XZR
- Intel x86 (not discussed much in this class)
 - 4 general purpose registers (eax, ebx, ecx, edx) 32 bits
 - Special registers: 3 pointer registers (si, di, ip), 4 segment (cs, ds, ss, es), 2 stack (sp, bp), status register (flags)
- LC2K (simple architecture made up for this class)
 - 8 registers, 32 bits each *(4 bytes)*



18

Agenda


- Computer Model and Binary
- ISAs
 - Registers
 - Control Flow**
 - Representing Different Values



20

How is Assembly Different from C/C++?

- C/C++: next line of code is executed until you get to:
 - function call
 - return statement
 - if statement or for/while loop
 - etc
- Assembly: a program counter (PC) keeps track of which memory address has the next instruction, gets incremented until
 - a "branch" or "jump" instruction *(to next instruction)*
 - Used to change control flow (more later)
 - This model is called a **von Neumann Architecture**



21

Traditional (von Neumann) Architecture

Here's the (endless) loop that hardware repeats forever:

- Fetch—get next instruction—use PC to find where it is in memory and place it in instruction register (IR)
 - PC is changed to "point" to the next instruction in the program
- Decode—control logic examines the contents of the IR to decide what instruction it should perform
- Execute—the outcome of the decoding process dictates
 - an arithmetic or logical operation on data
 - an access to data in the same memory as the instructions
 - OR a change to the contents of the PC




22

(Simplified) System Organization

Let's execute this short program (destination register listed first):


```
ADD X0, X1, X2
SUB X1, X2, X0
```



(Simplified) System Organization

Let's execute this short program (destination register listed first):


```
ADD X0, X1, X2
SUB X1, X2, X0
```



(Simplified) System Organization

Let's execute this short program
(destination register listed first):

ADD X0, X1, X2
SUB X1, X2, X0



Assembly Code – ARM Example

ARM v8 ISA

Poll: What are the final contents of X1, X2, and X3?

- What are the contents of the registers after executing the given assembly code (destination register is listed first in ARM)?

Program: `ADD X3, X1, X2`
`ADDI X3, X3, #3`
`SUB X2, X3, X1`

Initial register file:

X1	25
X2	-4
X3	57

ADDI means "add immediate", the last field is a literal value, not a register index.

(1) ADD X3, X1, X2

X1	25
X2	-4
X3	57

(2) ADDI X3, X3, #3

X1	25
X2	-4
X3	21

(3) SUB X2, X3, X1

X1	25
X2	-1
X3	24

Live Poll + Q&A: [slido.com #eeecs370](https://www.slido.com/join/eeecs370)

26

Agenda

- Computer Model and Binary
- ISAs
 - Registers
 - Control Flow
 - Representing Different Values

Different Data Types

- How does memory distinguish between different data types?
 - E.g. int, int *, char, float, double
- It doesn't! It's all just 0s and 1s!
- We'll see how to encode each of these later
- Exact length depends on architectures

How is Assembly Different from C/C++?

- No data types in assembly
- Everything is 0s and 1s: up to the programmer to interpret whether these bits should be interpreted as ints, bools, chars... or even instructions themselves!

```
char c = 'a';
c++; // c is now 'b'
```

// results in the same assembly as

```
int x = 97;
x++; // c is now 98
```

```
x = (int) c; // this instruction has no effect... why?
```

(这是 high level 语言的 compiler 限制的 compiler 的 implementation 要使这个高级语言合理 但是 assembly 没有那么多限制, 这跟 ISA 的 instructions 有关)

Minimum Datatype Sizes

Type	Minimum size (bits)
char	8
int	16
long int	32
float	32
double	64

Representing Values in Hardware

- Unsigned integers represented as we've seen
- Chars are represented as ASCII values
 - e.g. 'a' -> 97, 'b' -> 98, '#' -> 35
- What about negative numbers?
- Fractional numbers?

Negative Numbers

- There are many ways we could represent negative numbers
- Because it will eventually make our hardware simpler, the most common representation is 2's complement

(-748)

Hey, Good-Looking!

2

No, not 2's complement!

31

32

Two's Complement Representation

- Recall that 1101 in binary is 13 in decimal.

$$\begin{array}{cccc} 1 & 1 & 0 & 1 \\ 2^3 & 2^2 & 2^1 & 2^0 \end{array} = 8 + 4 + 1 = 13$$

- 2's complement numbers are very similar to unsigned binary numbers.
 - The only difference is that the first number is now negative.

$$\begin{array}{cccc} 1 & 1 & 0 & 1 \\ -2^3 & 2^2 & 2^1 & 2^0 \end{array} = -8 + 4 + 1 = -3$$

Fun with 2's Complement Numbers

- What is the range of representation of a 4-bit 2's complement number?
 - $[-8, 7]$ (corresponding to 1000 and 0111)
- What is the range of representation of an n-bit 2's complement number?
 - $[-2^{(n-1)}, 2^{(n-1)} - 1]$
- Useful trick: You can negate a 2's complement number by inverting all the bits and adding 1.

- 5 is represented as 0101
- Negate each bit: 1010
- Add 1: 1011 = -8 + 2 + 1 = -5

$$\begin{array}{r} 0100 (-5) \\ + 0101 (5) \\ \hline 1011 (-5) \end{array}$$




What about fractional numbers?

- One idea: fixed point notation
 - Have some bits represent numbers before decimal point, some bits represent numbers after decimal point
- Better idea: floating point notation
 - Inspired by scientific notation (e.g. 1.3×10^{-3})
 - Allows for larger range of numbers
 - We'll come back to this in a few lectures

Representing Instructions?

- Instructions, not just data, are stored in memory
- So, they must be expressible as numbers
- We'll look at how to encode instructions next time



Next Time

- Finish Up ISAs
- LC2K details
- Lingering questions / feedback? I'll include an anonymous form at the end of every lecture: <https://bit.ly/3oXr4Ah>



Addressing Modes

- Direct addressing
- Register indirect
- Base + displacement
- PC-relative



Direct Addressing

Example ISA (simplified)

- Consider this code:

```
const double PI = 3.14;  
double two_pi() {  
    return 2*PI;  
}
```

Not practical in modern ISAs...
if we have 32 bit instructions
and 32 bit addresses, the entire
instruction is the address!

- When we load PI, it's ALWAYS the same address
 - If the ISA supports it, we can just hardcode that address in the instruction
- Like register addressing
 - Specify address as immediate constant

```
load r1, mem[1500] ; r1 ← contents of location 1500  
jump mem[3000] ; jump to address 3000
```

- Useful for addressing locations that don't change during execution
 - Branch target addresses
 - Global/static variable locations



Register indirect

Example ISA (simplified)

- Consider this code:

```
int my_arr[2] = {6666, 7777};  
int* ptr = &my_arr[0];  
for(int i=0; i<2; i++) {  
    int x = *ptr;  
    ptr++;  
}
```

(寄存器间接寻址)
使用寄存器保存地址

- Everytime we load into x, it's a different address
- But the address is always stored in another variable
- If ISA supports it, we could use a load like this
 - load r1, mem[r2]




Register indirect

- Consider this code:

```
int my_arr[2] = {6666, 7777};
int* ptr = &my_arr[0];
for(int i=0; i<2; i++) {
    int x = *ptr;
    ptr++;
}
```

```
load r1, mem[ r2 ]
add r2, r2, #4
load r1, mem[ r2 ]
```




Example ISA (simplified)

Register indirect

- Consider this code:

```
int my_arr[2] = {6666, 7777};
int* ptr = &my_arr[0];
for(int i=0; i<2; i++) {
    int x = *ptr;
    ptr++;
}
```

```
load r1, mem[ r2 ]
add r2, r2, #4
load r1, mem[ r2 ]
```




Example ISA (simplified)

Register indirect

- Consider this code:

```
int my_arr[2] = {6666, 7777};
int* ptr = &my_arr[0];
for(int i=0; i<2; i++) {
    int x = *ptr;
    ptr++;
}
```

```
load r1, mem[ r2 ]
add r2, r2, #4
load r1, mem[ r2 ]
```



This is better, but we can be more general


Example ISA (simplified)

Base + Displacement

- Consider this code:

```
struct My_Struct {
    int tot;
    //...
    int val;
};
My_Struct a;
//...
a.tot += a.val;
```

load r1, mem[r2 + 32]



- If a register holds the starting address of "a" ...
 - Then the specific values needed are a slight offset
- Base + Displacement
 - reg value + immed


Very general, most common addressing mode today

Class Problem

- a. What are the contents of register/memory after executing the following instructions

```
r2 = load mem[r3]
r3 = load mem[r2+4]
store mem[r2+8], r3
```

Q: What are the contents of register / memory?



PC-relative addressing

- Relevant for P1.a!

- Variant on base + displacement
- Remember PC is "Program Counter", keeps track of which line (memory address) of the program we're at
- PC register is base, longer displacement possible since PC is assumed implicitly (more bits available)
 - Used for branch instructions
 - jump [-8]; jump back 2 instructions (32-bit instructions)

PC是一个register, 专用于存放
在遇到哪条指令

ISA Types

Reduced Instruction Set Computing (RISC)

- Fewer, simpler instructions
- Encoding of instructions are usually the same size
- Simpler hardware
- Program is larger, more tedious to write by hand
- E.g. LC2K, RISC-V, (ARM) kinda
- More popular now

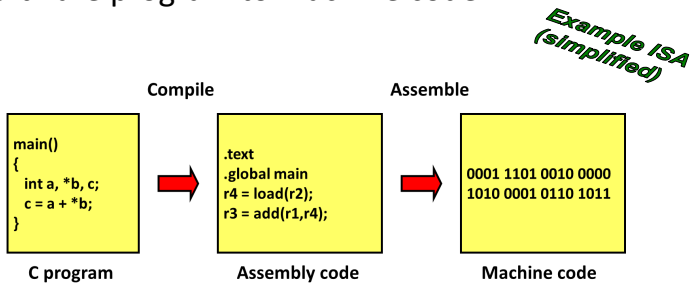
Complex Instruction Set Computing (CISC)

- More, complex instructions
- Encoding of instructions are different sizes
- More complex hardware
- Short, expressive programs, easier to write by hand
- E.g. x86
- Less popular now

Encoding Instructions


- So binary numbers can represent signed and unsigned numbers, chars, and fractional numbers
- But they must also represent instructions themselves!
 - After all, memory is just a collection of 1s and 0s
- We need a way of **encoding** instructions in order to store them in memory

Software program to machine code



Assembly Instruction Encoding

- Since the EDSAC (1949) almost all computers stored program instructions the same way they store data.
- Each instruction is encoded as a number



- This is the number stored in memory (in binary)! **Poll:** How many different "operation codes" could be supported by this ISA? How many registers?

Operating on Binary Values

- All values are stored in binary, even when you specify the number in decimal
- It is often convenient to treat values as sequences of bits, rather than values
 - You will need to do this in P1a
- C provides "bitwise operators" to do this
 - Shift ("<<" and ">>")
 - Bitwise boolean ("&", "|", "^", and "~")

Shift Operators

- Shift a value x bits to the left via "<<"
- Inserts "x" zeros to the right (least significant)
- E.g.

```
int a = 60;
int s = a << 2;
```

Shift Operators

- Shift a value x bits to the left via "<<"
- Inserts "x" zeros to the right (least significant)
- E.g.


```
int a = 60; // 0b0011_1100
int s = a << 2; // 0b1111_0000
```
- "a" is still 60, "s" is 240
- Same idea for ">>", but to the right

shifting x to the left in decimal → multiplying by 10^x
shifting x to the left in binary → multiplying by 2^x

Bitwise operations

- Bitwise operations apply a Boolean operation on each bit of a value (or each pair of bits across two values)

```
int a = 60; // 0b0011_1100
int b = 13; // 0b0000_1101
int o = a | b; // 0b0011_1101
```

- "a" and "b" are the same, "o" is 61
- & – and | – or ^ – xor ~ – not
- Very different from Boolean &&, ||, etc
 - Why?